

Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting


Cluster Architecture

- ❑ Kubernetes Architecture
- ❑ ETCD For Beginners
- ❑ ETCD in Kubernetes
- ❑ Kube-API Server
- ❑ Controller Managers
- ❑ Kube Scheduler
- ❑ Kubelet
- ❑ Kube Proxy



KUBERNETES ARCHITECTURE







Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers

kube-apiserver



Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers

kubelet

Controller-
Manager


D
CLUSTER


e-scheduler

Kube-proxy

Container Runtime Engine

Run containers





Kubernetes Architecture



Master

Manage, Plan, Schedule, Monitor
Nodes

ETCD
CLUSTER

kube-
apiserver

Kube
Controller
Manager

kube-scheduler



Worker Nodes

Host Application as Containers

kubelet

Kube-proxy

Container Runtime Engine

Run containers



rkt

kubelet

Kube-proxy

Container Runtime Engine

Run containers



rkt

ETCD FOR BEGINNERS



Objectives

- What is ETCD?
- What is a Key-Value Store?
- How to get started quickly?
- How to operate ETCD?
- What is a distributed system?
- How ETCD Operates
- RAFT Protocol
- Best practices on number of nodes



**ETCD is a distributed
reliable key-value store
that is Simple, Secure &
Fast**

key-value store

Tabular/Relational Databases

Name	Age	Location	Salary	Grade
John Doe	45	New York	5000	
Dave Smith	34	New York	4000	
Aryan Kumar	10	New York		A
Lauren Rob	13	Bangalore		C
Lily Oliver	15	Bangalore		B

key-value store

Key	Value
Name	John Doe
Age	45
Location	New York
Salary	5000

Key	Value
Name	Dave Smith
Age	34
Location	New York
Salary	4000
Organization	ACME

Key	Value
Name	Aryan Kumar
Age	10
Location	New York
Grade	A

Key	Value
Name	Lauren Rob
Age	13
Location	Bangalore
Grade	C

Key	Value
Name	Lily Oliver
Age	15
Location	Bangalore
Grade	B

key-value store

```
{  
  "name": "John Doe",  
  "age": 45,  
  "location": "New York",  
  "salary": 5000  
}
```

```
{  
  "name": "Dave Smith",  
  "age": 34,  
  "location": "New York",  
  "salary": 4000,  
  "organization": "ACME"  
}
```

```
{  
  "name": "Aryan Kumar",  
  "age": 10,  
  "location": "New York",  
  "Grade": "A"  
}
```

```
{  
  "name": "Lily Oliver",  
  "age": 15,  
  "location": "Bangalore",  
  "Grade": "B"  
}
```

```
{  
  "name": "Lauren Rob",  
  "age": 13,  
  "location": "Bangalore",  
  "Grade": "C"  
}
```


Install ETCD

1. Download Binaries

```
curl -L https://github.com/etcd-io/etcd/releases/download/v3.3.11/etcd-v3.3.11-linux-amd64.tar.gz -o etcd-v3.3.11-linux-amd64.tar.gz
```

2. Extract

```
tar xzvf etcd-v3.3.11-linux-amd64.tar.gz
```

3. Run ETCD Service

```
./etcd
```

Operate ETCD

3. Run ETCD Service

```
./etcd
```

```
./etcdctl set key1 value1
```

```
./etcdctl get key1
```

```
value1
```

```
./etcdctl
```

NAME:

etcdctl - A simple command line client for etcd.

COMMANDS:

backup	backup an etcd directory
cluster-health	check the health of the etcd cluster
mk	make a new key with a given value
mkdir	make a new directory
rm	remove a key or a directory
rmdir	removes the key if it is an empty directory or a key-value pair
get	retrieve the value of a key



Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

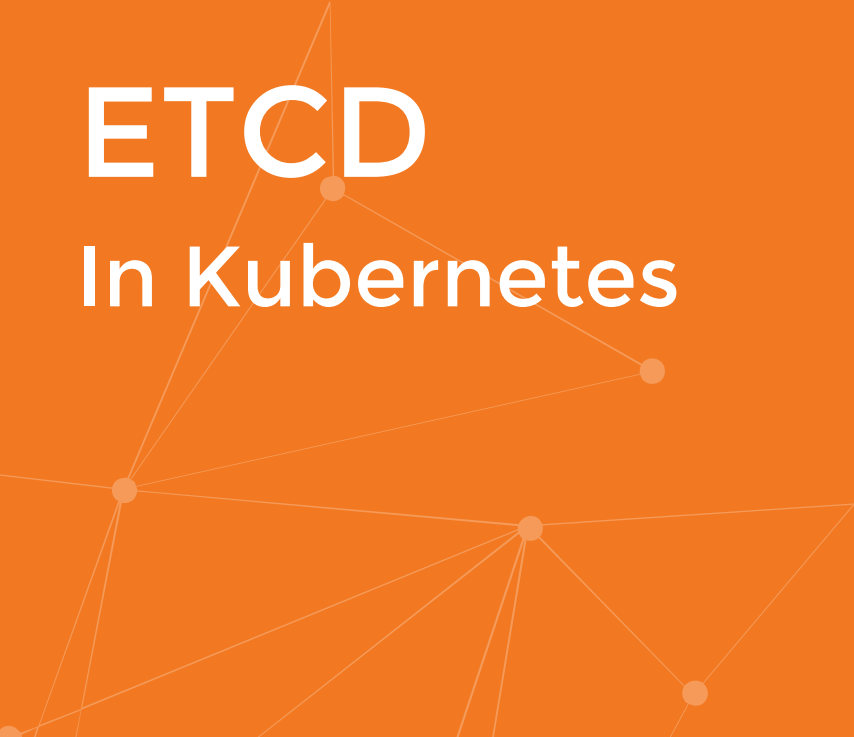
Networking

Installation, Configuration & Validation

Troubleshooting

ETCD

In Kubernetes






Master

Manage, Plan, Schedule, Monitor
Nodes

- Nodes
- PODs
- Configs
- Secrets
- Accounts
- Roles
- Bindings
- Others

**ETCD
CLUSTER**



Setup - Manual

```
▶ wget -q --https-only \  
  "https://github.com/coreos/etcd/releases/download/v3.3.9/etcd-v3.3.9-linux-amd64.tar.gz"
```

etcd.service

```
ExecStart=/usr/local/bin/etcd \  
  --name ${ETCD_NAME} \  
  --cert-file=/etc/etcd/kubernetes.pem \  
  --key-file=/etc/etcd/kubernetes-key.pem \  
  --peer-cert-file=/etc/etcd/kubernetes.pem \  
  --peer-key-file=/etc/etcd/kubernetes-key.pem \  
  --trusted-ca-file=/etc/etcd/ca.pem \  
  --peer-trusted-ca-file=/etc/etcd/ca.pem \  
  --peer-client-cert-auth \  
  --client-cert-auth \  
  --initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \  
  --listen-peer-urls https://${INTERNAL_IP}:2380 \  
  --listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \  
  --advertise-client-urls https://${INTERNAL_IP}:2379 \  
  --initial-cluster-token etcd-cluster-0 \  
  --initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \  
  --initial-cluster-state new \  
  --data-dir=/var/lib/etcd
```

Setup - kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-prwv1	1/1	Running	0	1h
kube-system	coredns-78fcd6894-vqd9w	1/1	Running	0	1h
kube-system	etcd-master	1/1	Running	0	1h
kube-system	kube-apiserver-master	1/1	Running	0	1h
kube-system	kube-controller-manager-master	1/1	Running	0	1h
kube-system	kube-proxy-f6k26	1/1	Running	0	1h
kube-system	kube-proxy-hnzsx	1/1	Running	0	1h
kube-system	kube-scheduler-master	1/1	Running	0	1h
kube-system	weave-net-924k8	2/2	Running	1	1h
kube-system	weave-net-hzfcz	2/2	Running	1	1h

```
kubectl exec etcd-master -n kube-system etcdctl get / --prefix --keys-only
```

```
/registry/apiregistration.k8s.io/apiservices/v1.  
/registry/apiregistration.k8s.io/apiservices/v1.apps  
/registry/apiregistration.k8s.io/apiservices/v1.authentication.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.autoscaling  
/registry/apiregistration.k8s.io/apiservices/v1.batch  
/registry/apiregistration.k8s.io/apiservices/v1.networking.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.rbac.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.storage.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1beta1.admissionregistration.k8s.io
```

Run inside the etcd-
master POD

Explore ETCD

```
▶ kubectl exec etcd-master -n kube-system etcdctl get / --prefix --keys-only  
  
/registry/apiregistration.k8s.io/apiservices/v1.  
/registry/apiregistration.k8s.io/apiservices/v1.apps  
/registry/apiregistration.k8s.io/apiservices/v1.authentication.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.autoscaling  
/registry/apiregistration.k8s.io/apiservices/v1.batch  
/registry/apiregistration.k8s.io/apiservices/v1.networking.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.rbac.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.storage.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1beta1.admissionregistration.k8s.io
```

Run inside the etcd-
master POD

Registry

minions

pods

replicasets

deployments

roles

secrets

ETCD in HA Environment



etcd.service

```
ExecStart=/usr/local/bin/etcd \\  
  --name ${ETCD_NAME} \\  
  --cert-file=/etc/etcd/kubernetes.pem \\  
  --key-file=/etc/etcd/kubernetes-key.pem \\  
  --peer-cert-file=/etc/etcd/kubernetes.pem \\  
  --peer-key-file=/etc/etcd/kubernetes-key.pem \\  
  --trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-client-cert-auth \\  
  --client-cert-auth \\  
  --initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \\  
  --advertise-client-urls https://${INTERNAL_IP}:2379 \\  
  --initial-cluster-token etcd-cluster-0 \\  
  --initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \\  
  --initial-cluster-state new \\  
  --data-dir=/var/lib/etcd
```


Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

kube-api server



kube-apiserver



Master

Manage, Plan, Schedule, Monitor
Nodes

kubelet




Worker Nodes

Host Application as Containers

Controller-
Manager

D
CLUSTER

e-scheduler



Kubernetes Architecture



```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	20m	v1.11.3
node01	Ready	<none>	20m	v1.11.3

1. Authenticate User

kube-apiserver

2. Validate Request

3. Retrieve data

Controller-
manager

ETCD
CLUSTER

kube-scheduler



Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers

kubelet

Container Runtime Engine

Run containers



Worker Nodes

Host Application as Containers

kubelet

Container Runtime Engine

Run containers



Kubernetes Architecture



```
curl -X POST /api/v1/namespaces/default/pods ...[other]  
Pod created!
```

kube-apiserver

Controller-
manager

ETCD
CLUSTER

kube-scheduler



Master

Manage, Plan, Schedule, Monitor
Nodes

kubelet

Worker Nodes

Host Application as Containers

Container Runtime Engine

Run containers



Worker Nodes

Host Application as Containers

Container Runtime Engine

Run containers



1. Authenticate User

2. Validate Request

3. Retrieve data

4. Update ETCD

5. Scheduler

6. Kubelet

Kube-api Server

1. Authenticate User

2. Validate Request

3. Retrieve data

4. Update ETCD

5. Scheduler

6. Kubelet

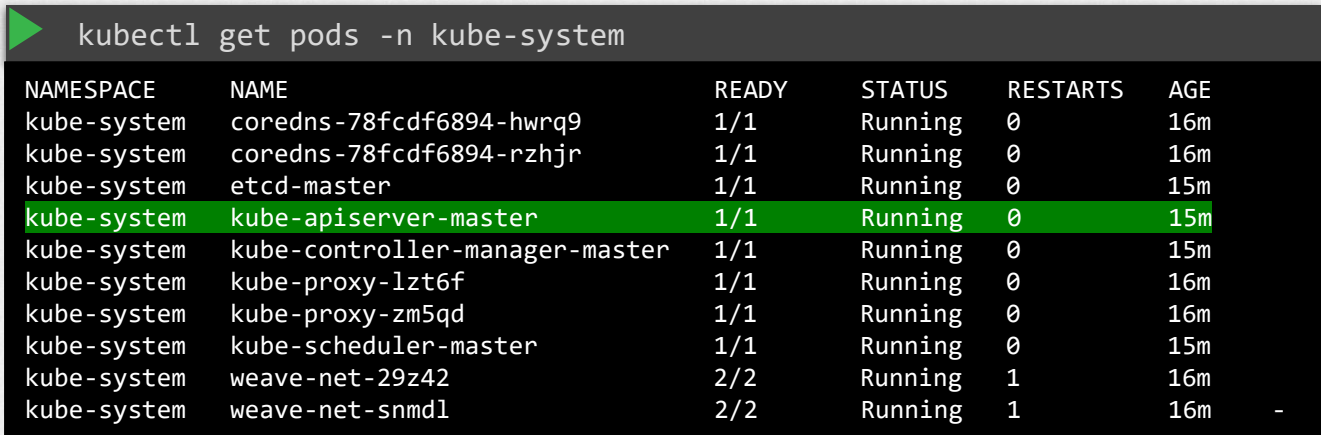
Installing kube-api server

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-apiserver
```

kube-apiserver.service

```
ExecStart=/usr/local/bin/kube-apiserver \\  
  --advertise-address=${INTERNAL_IP} \\  
  --allow-privileged=true \\  
  --apiserver-count=3 \\  
  --authorization-mode=Node,RBAC \\  
  --bind-address=0.0.0.0 \\  
  --enable-admission-  
plugins=Initializers,NamespaceLifecycle,NodeRestriction,LimitRanger,ServiceAccount,DefaultStorageClass,ResourceQuota \\  
  --enable-swagger-ui=true \\  
  --etcd-servers=https://127.0.0.1:2379 \\  
  --event-ttl=1h \\  
  --experimental-encryption-provider-config=/var/lib/kubernetes/encryption-config.yaml \\  
  --runtime-config=api/all \\  
  --service-account-key-file=/var/lib/kubernetes/service-account.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --service-node-port-range=30000-32767 \\  
  --v=2
```

View api-server - kubeadm



```
▶ kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

View api-server options - kubeadm

```
cat /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
spec:
  containers:
  - command:
    - kube-apiserver
    - --authorization-mode=Node,RBAC
    - --advertise-address=172.17.0.32
    - --allow-privileged=true
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --disable-admission-plugins=PersistentVolumeLabel
    - --enable-admission-plugins=NodeRestriction
    - --enable-bootstrap-token-auth=true
    - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
    - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
    - --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key
    - --etcd-servers=https://127.0.0.1:2379
    - --insecure-port=0
    - --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt
    - --kubelet-client-key=/etc/kubernetes/pki/apiserver-kubelet-client.key
    - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
    - --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt
    - --proxy-client-key-file=/etc/kubernetes/pki/front-proxy-client.key
    - --requestheader-allowed-names=front-proxy-client
    - --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt
    - --requestheader-extra-headers-prefix=X-Remote-Extra-
    - --requestheader-group-headers=X-Remote-Group
    - --requestheader-username-headers=X-Remote-User
```

View api-server options

```
cat /etc/systemd/system/kube-apiserver.service
```

```
[Service]
ExecStart=/usr/local/bin/kube-apiserver \\\
  --advertise-address=${INTERNAL_IP} \\\
  --allow-privileged=true \\\
  --apiserver-count=3 \\\
  --audit-log-maxage=30 \\\
  --audit-log-maxbackup=3 \\\
  --audit-log-maxsize=100 \\\
  --audit-log-path=/var/log/audit.log \\\
  --authorization-mode=Node,RBAC \\\
  --bind-address=0.0.0.0 \\\
  --client-ca-file=/var/lib/kubernetes/ca.pem \\\
  --enable-admission-
plugins=Initializers,NamespaceLifecycle,NodeRestriction,LimitRanger,ServiceAccount,DefaultStorageClass,ResourceQuota \\\
  --enable-swagger-ui=true \\\
  --etcd-cafile=/var/lib/kubernetes/ca.pem \\\
  --etcd-certfile=/var/lib/kubernetes/kubernetes.pem \\\
  --etcd-keyfile=/var/lib/kubernetes/kubernetes-key.pem \\\
  --etcd-
servers=https://10.240.0.10:2379,https://10.240.0.11:2379,https://10.240.0.12:2379 \\\
  --event-ttl=1h \\\
  --experimental-encryption-provider-config=/var/lib/kubernetes/encryption-config.yaml \\\
  --kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \\\
  --kubelet-client-certificate=/var/lib/kubernetes/kubernetes.pem \\\
```

View api-server options

```
▶ ps -aux | grep kube-apiserver

root      2348   3.3 15.4 399040 315604 ?        Ssl  15:46   1:22 kube-apiserver --authorization-mode=Node,RBAC --
advertise-address=172.17.0.32 --allow-privileged=true --client-ca-file=/etc/kubernetes/pki/ca.crt --disable-
admission-plugins=PersistentVolumeLabel --enable-admission-plugins=NodeRestriction --enable-bootstrap-token-
auth=true --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-
client.crt --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key --etcd-servers=https://127.0.0.1:2379 --
insecure-port=0 --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt --kubelet-client-
key=/etc/kubernetes/pki/apiserver-kubelet-client.key --kubelet-preferred-address-
types=InternalIP,ExternalIP,Hostname --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt --proxy-
client-key-file=/etc/kubernetes/pki/front-proxy-client.key --requestheader-allowed-names=front-proxy-client --
requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt --requestheader-extra-headers-prefix=X-Remote-
Extra- --requestheader-group-headers=X-Remote-Group --requestheader-username-headers=X-Remote-User --secure-
port=6443 --service-account-key-file=/etc/kubernetes/pki/sa.pub --service-cluster-ip-range=10.96.0.0/12 --tls-
cert-file=/etc/kubernetes/pki/apiserver.crt --tls-private-key-file=/etc/kubernetes/pki/apiserver.key
```


Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

Kube Controller Manager





Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers

Controller-Manager


ETCD
CLUSTER

kube-scheduler


Watch Status

Remediate Situation

Controller



kube-apiserver



Watch Status

Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s


POD Eviction Timeout = 5m

```
► kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0

UNREACHABLE

Controller



Watch Status

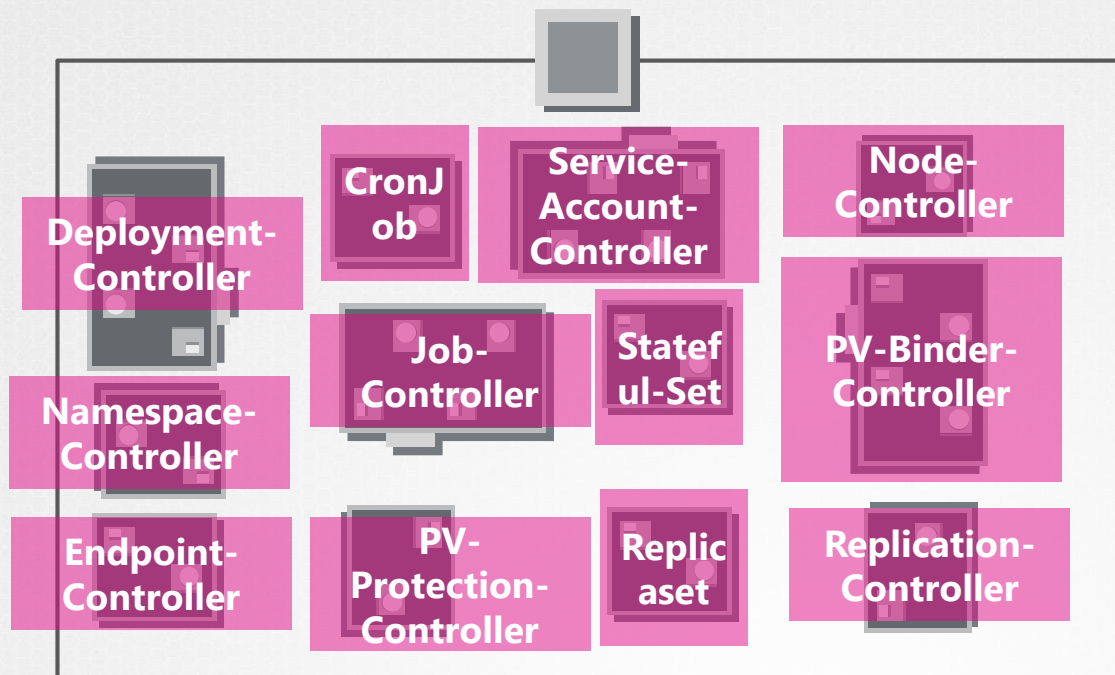
Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

Controller



Watch Status

Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

Controller



Watch Status

Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

Installing kube-controller-manager

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-controller-manager
```

kube-controller-manager.service

```
ExecStart=/usr/local/bin/kube-controller-manager \\  
  --address=0.0.0.0 \\  
  --cluster-cidr=10.200.0.0/16 \\  
  --cluster-name=kubernetes \\  
  --cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \\  
  --cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \\  
  --kubeconfig=/var/lib/kubernetes/kube-controller-manager.kubeconfig \\  
  --leader-elect=true \\  
  --root-ca-file=/var/lib/kubernetes/ca.pem \\  
  --service-account-private-key-file=/var/lib/kubernetes/service-account-key.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --use-service-account-credentials=true \\  
  --v=2
```

```
N --node-monitor-period=5s
```

```
--controllers stringSlice      Default: [*]
```

A list of controllers to enable. '*' enables all on-by-default controllers, 'foo' enables the controller named 'foo', '-foo' disables the controller named 'foo'.

All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving, csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector, horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder, persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller,

```
d=40s
```

```
s
```


Installing kube-controller-manager

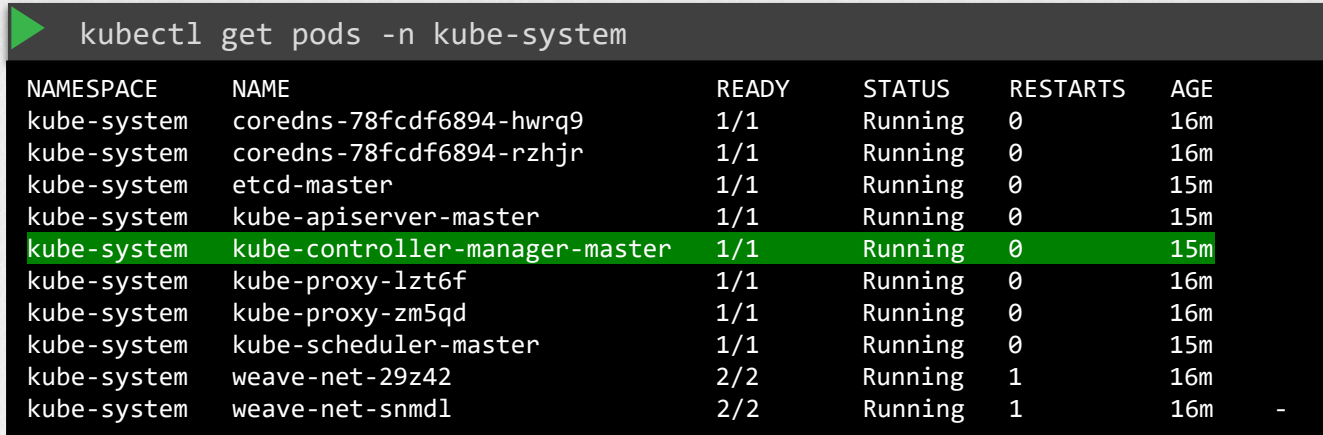
```
--controllers stringSlice      Default: [*]
```

A list of controllers to enable. '*' enables all on-by-default controllers, 'foo' enables the controller named 'foo', '-foo' disables the controller named 'foo'.

All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving, csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector, horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder, persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller, resourcequota, root-ca-cert-publisher, route, service, serviceaccount, serviceaccount-token, statefulset, tokencleaner, ttl, ttl-after-finished

Disabled-by-default controllers: bootstrapsigner, tokencleaner

View kube-controller-manager - kubeadm



```
▶ kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

View kube-controller-manager options

- kubeadm

```
▶ cat /etc/kubernetes/manifests/kube-controller-manager.yaml
```

```
spec:
  containers:
  - command:
    - kube-controller-manager
    - --address=127.0.0.1
    - --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt
    - --cluster-signing-key-file=/etc/kubernetes/pki/ca.key
    - --controllers=*,bootstrapsigner,tokencleaner
    - --kubeconfig=/etc/kubernetes/controller-manager.conf
    - --leader-elect=true
    - --root-ca-file=/etc/kubernetes/pki/ca.crt
    - --service-account-private-key-file=/etc/kubernetes/pki/sa.key
    - --use-service-account-credentials=true
```

View controller-manager options

```
cat /etc/systemd/system/kube-controller-manager.service
```

```
[Service]
ExecStart=/usr/local/bin/kube-controller-manager \\\
  --address=0.0.0.0 \\\
  --cluster-cidr=10.200.0.0/16 \\\
  --cluster-name=kubernetes \\\
  --cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \\\
  --cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \\\
  --kubeconfig=/var/lib/kubernetes/kube-controller-manager.kubeconfig \\\
  --leader-elect=true \\\
  --root-ca-file=/var/lib/kubernetes/ca.pem \\\
  --service-account-private-key-file=/var/lib/kubernetes/service-account-key.pem \\\
  --service-cluster-ip-range=10.32.0.0/24 \\\
  --use-service-account-credentials=true \\\
  --v=2
Restart=on-failure
RestartSec=5
```

View controller-manager options

```
▶ ps -aux | grep kube-controller-manager  
  
root      1994   2.7   5.1 154360 105024 ?        Ssl  06:45   1:25 kube-controller-manager --  
address=127.0.0.1 --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt --cluster-signing-  
key-file=/etc/kubernetes/pki/ca.key --controllers=*,bootstrapsigner,tokencleaner --  
kubeconfig=/etc/kubernetes/controller-manager.conf --leader-elect=true --root-ca-  
file=/etc/kubernetes/pki/ca.crt --service-account-private-key-file=/etc/kubernetes/pki/sa.key  
--use-service-account-credentials=true
```


Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

Kube Scheduler





Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes


Host Application as Containers

Controller-
Manager


ETCD
CLUSTER




Kube-Scheduler




Kube-Scheduler



Kube-Scheduler





Kube-Scheduler



1. Filter Nodes

2. Rank Nodes



| More Later...

- Resource Requirements and Limits
- Taints and Tolerations
- Node Selectors/Affinity

Course Objectives

Scheduling

- ☐ Labels & Selectors
- ☐ Resource Limits
- ☐ Manual Scheduling
- ☐ Daemon Sets
- ☐ Multiple Schedulers
- ☐ Scheduler Events
- ☐ Configure Kubernetes Scheduler

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Troubleshooting

Installing kube-scheduler

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-scheduler
```

kube-scheduler.service

```
ExecStart=/usr/local/bin/kube-scheduler \\  
--config=/etc/kubernetes/config/kube-scheduler.yaml \\  
--v=2
```


View kube-scheduler options - kubeadm

```
▶ cat /etc/kubernetes/manifests/kube-scheduler.yaml
```

```
spec:
  containers:
  - command:
    - kube-scheduler
    - --address=127.0.0.1
    - --kubeconfig=/etc/kubernetes/scheduler.conf
    - --leader-elect=true
```

View kube-scheduler options

```
▶ ps -aux | grep kube-scheduler  
root      2477  0.8  1.6 48524 34044 ?        Ssl  17:31   0:08 kube-scheduler --  
address=127.0.0.1 --kubeconfig=/etc/kubernetes/scheduler.conf --leader-elect=true
```


Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

Kubelet



kube-apiserver



Master

Manage, Plan, Schedule, Monitor
Nodes

kubelet



Worker Nodes

Host Application as Containers

Controller-
Manager

D
CLUSTER

e-scheduler



docker

Kubernetes Architecture



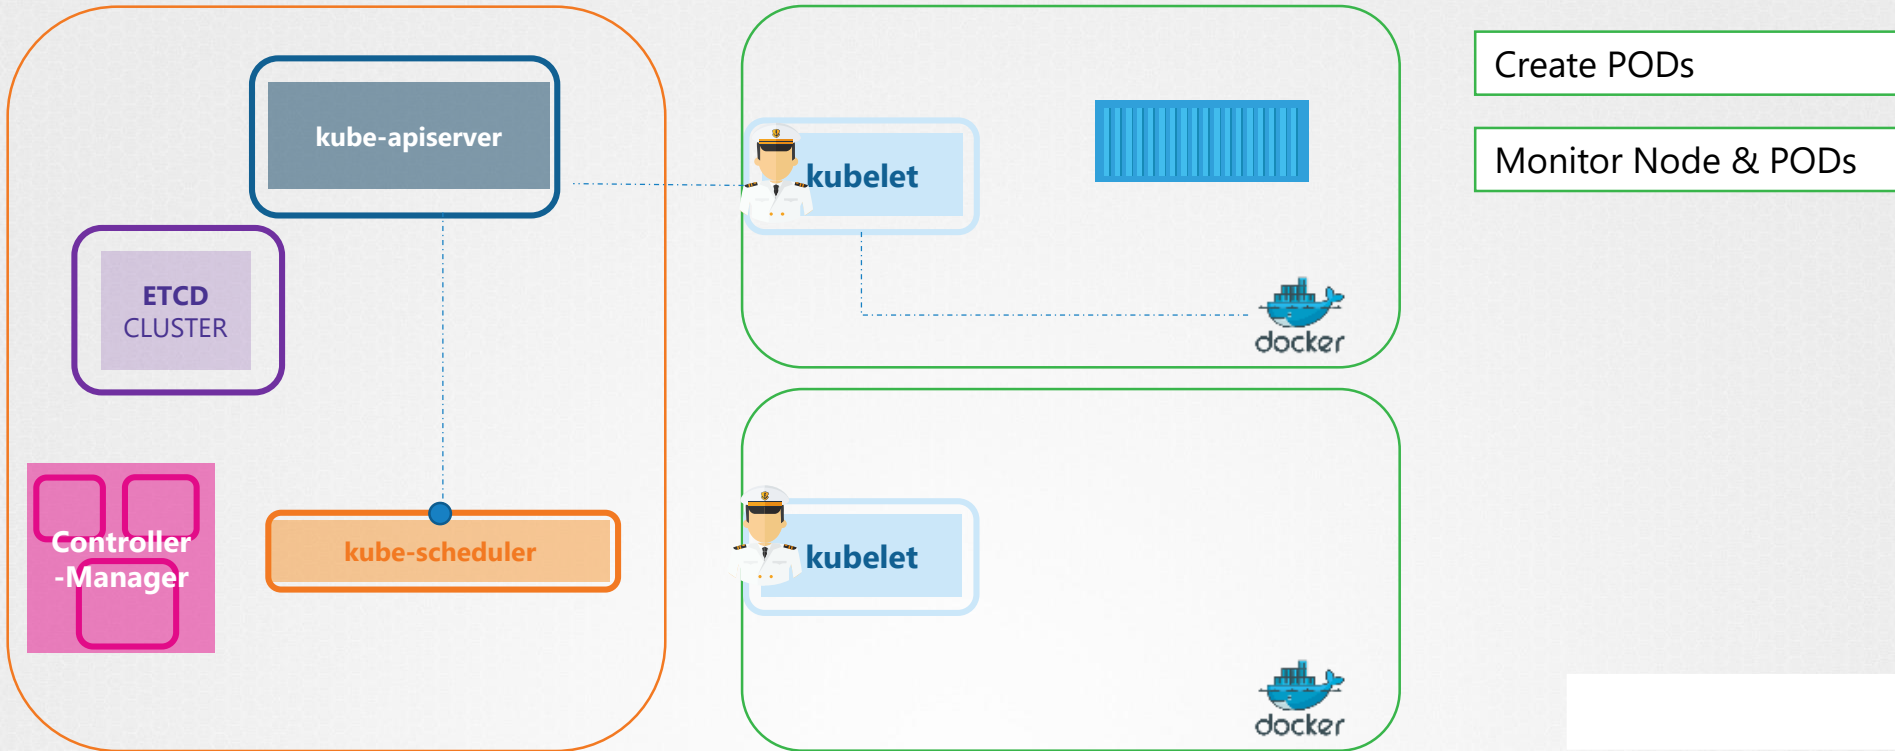
Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers



Installing kubelet

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kubelet
```

kubelet.service

```
ExecStart=/usr/local/bin/kubelet \\  
  --config=/var/lib/kubelet/kubelet-config.yaml \\  
  --container-runtime=remote \\  
  --container-runtime-endpoint=unix:///var/run/containerd/containerd.sock \\  
  --image-pull-progress-deadline=2m \\  
  --kubeconfig=/var/lib/kubelet/kubeconfig \\  
  --network-plugin=cni \\  
  --register-node=true \\  
  --v=2
```



Kubeadm does not deploy
Kubelets

View kubelet options

```
▶ ps -aux | grep kubelet  
  
root      2095  1.8  2.4 960676 98788 ?        Ssl  02:32   0:36 /usr/bin/kubelet --bootstrap-  
kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --  
config=/var/lib/kubelet/config.yaml --cgroup-driver=cgroupfs --cni-bin-dir=/opt/cni/bin --cni-  
conf-dir=/etc/cni/net.d --network-plugin=cni
```


Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

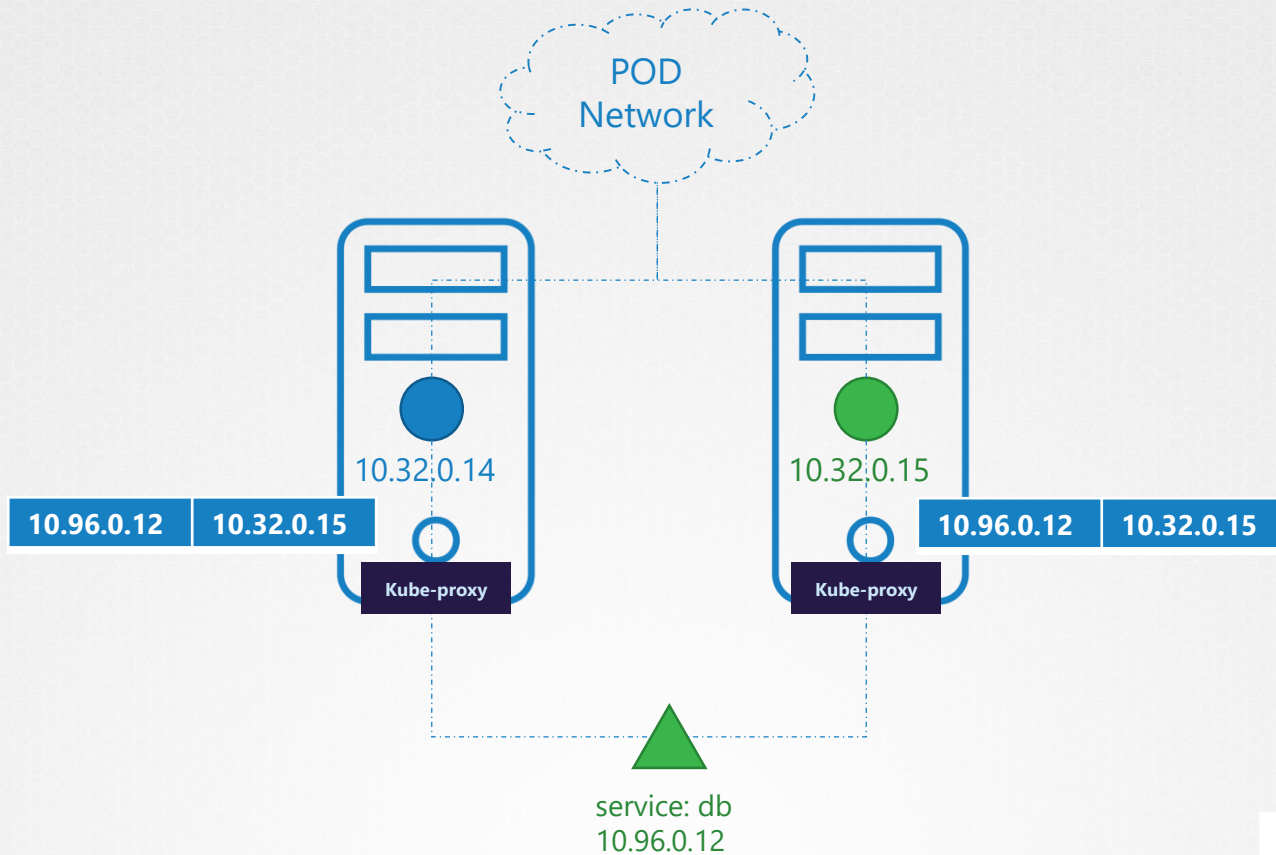
Installation, Configuration & Validation

Troubleshooting

Kube-proxy

An abstract network diagram consisting of several orange dots connected by thin orange lines, forming a web-like structure. The dots are positioned at various points, with some acting as central hubs connected to multiple other nodes.

Kube-proxy



Installing kube-proxy

```
▶ wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-proxy
```

kube-proxy.service

```
ExecStart=/usr/local/bin/kube-proxy \\  
  --config=/var/lib/kube-proxy/kube-proxy-config.yaml  
Restart=on-failure  
RestartSec=5
```

View kube-proxy - kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

```
kubectl get daemonset -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-proxy	2	2	2	2	2	beta.kubernetes.io/arch=amd64	1h