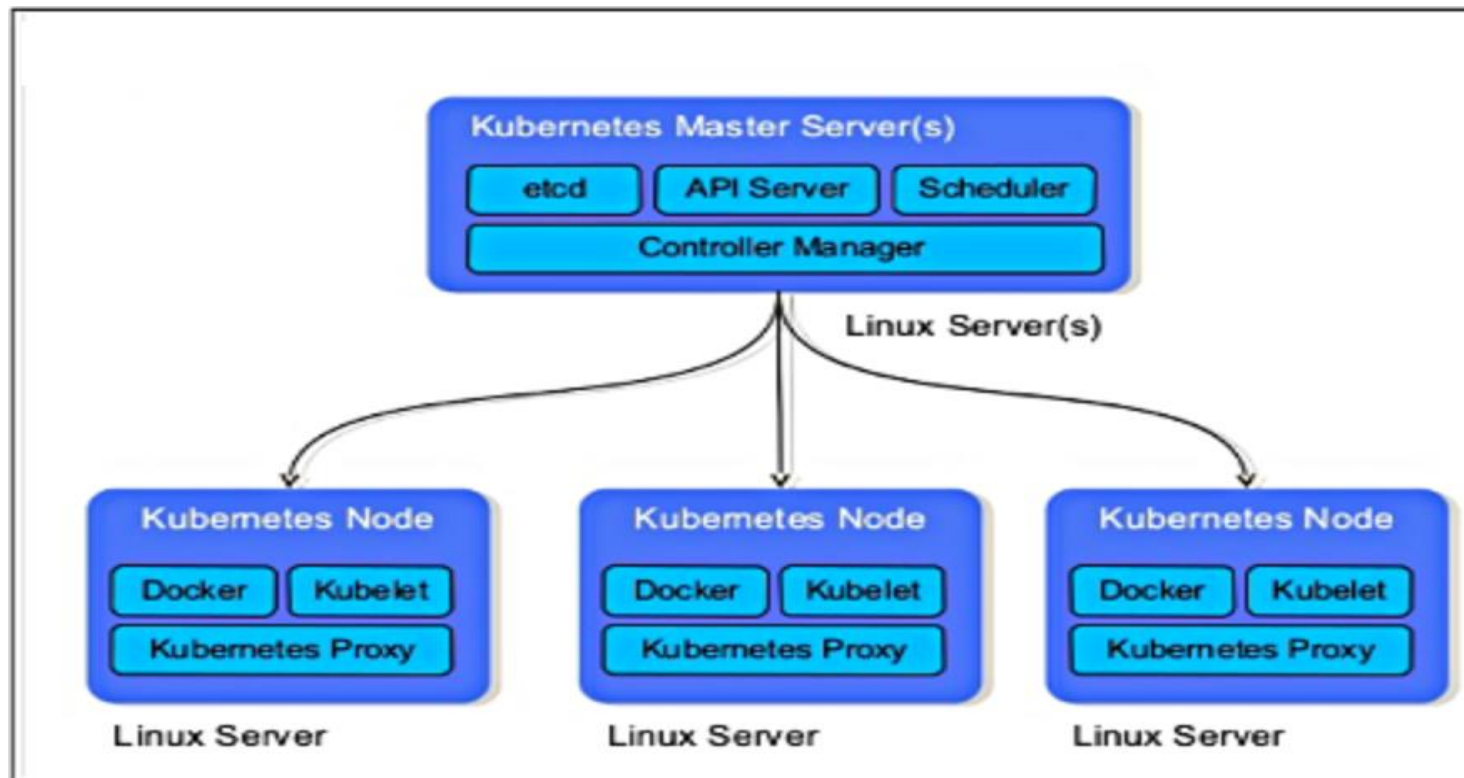


Docker Container Clustering using Kubernetes

- Kubernetes is a container management technology developed in Google lab
- Kubernetes is an open source container management tool hosted by Cloud Native Computing Foundation (CNCF)
- Kubernetes also called "K8s" is an open-source system for automating deployment, management and scaling of containerized applications
- It supports a range of container tools, including Docker
- Kubernetes v1.0 was released on July 21, 2015

- Kubernetes comes with a capability of automating deployment, scaling of application, and operations of application containers across clusters
- Following are some of the important features of Kubernetes
 - Continuous development, integration and deployment
 - Containerized infrastructure
 - Application-centric management
 - Auto-scalable infrastructure
 - Environment consistency across development testing and production
 - Loosely coupled infrastructure, where each component can act as a separate unit
 - Higher density of resource utilization
 - Predictable infrastructure which is going to be created

Kubernetes Architecture



Kubernetes – Master Machine Components

- Following are the components of Kubernetes Master Machine
 - etcd
 - Stores the configuration information which can be used by each of the nodes in the cluster
 - It is a high availability key value store that can be distributed among multiple nodes
 - It is accessible only by Kubernetes API server as it may have some sensitive information
 - It is a distributed key value Store which is accessible to all
 - API Server
 - Kubernetes is an API server which provides all the operation on cluster using the API
 - API server implements an interface, which means different tools and libraries can readily communicate with it
 - **Kubeconfig** is a package along with the server side tools that can be used for communication
 - It exposes Kubernetes API

Kubernetes – Master Machine Components Contd.

- **Controller Manager**
 - This component is responsible for most of the controllers that regulate the state of cluster and perform a task
 - It can be considered as a daemon which runs in a nonterminating loop and is responsible for collecting and sending information to API server
 - It works toward getting the shared state of cluster and then make changes to bring the current status of the cluster to the desired state.
 - The key controllers are replication controller, endpoint controller, namespace controller, and service account controller.
 - The controller manager runs different kind of controllers to handle nodes, endpoints, etc.
- **Scheduler**
 - This is one of the key components of Kubernetes master
 - It is a service in master responsible for distributing the workload
 - It is responsible for tracking utilization of working load on cluster nodes and then placing the workload on which resources are available and accept the workload
 - This is the mechanism responsible for allocating pods to available nodes
 - The scheduler is responsible for workload utilization and allocating pod to new node

Kubernetes – Node Components

- Following are the key components of Node server which are necessary to communicate with Kubernetes master
 - Docker
 - The first requirement of each node is Docker which helps in running the encapsulated application containers in a relatively isolated but lightweight operating environment
 - Kubelet Service
 - This is a small service in each node responsible for relaying information to and from control plane service
 - It interacts with **etcd** store to read configuration details and right values
 - This communicates with the master component to receive commands and work
 - The **kubelet** process then assumes responsibility for maintaining the state of work and the node server
 - It manages network rules, port forwarding, etc
 - Kubernetes Proxy Service
 - This is a proxy service which runs on each node and helps in making services available to the external host
 - It helps in forwarding the request to correct containers and is capable of performing primitive load balancing
 - It makes sure that the networking environment is predictable and accessible and at the same time it is isolated as well
 - It manages pods on node, volumes, secrets, creating new containers' health check up, etc

Kubernetes – Installation

Google Container Services

- The Google Cloud Platform offers a hosted Kubernetes-as-a-Service called Google Container Engine (GKE). To get started with GKE, you need a Google Cloud Platform account with billing enabled and the gcloud tool installed
- Once you have gcloud installed, first set a default zone:
`gcloud config set compute/zone us-west1-a`
- Then you can create a cluster:
`gcloud container clusters create kuar-cluster`
- When the cluster is ready you can get credentials for the cluster using:
`gcloud auth application-default login`
- At this point, you should have a cluster configured and ready to go

Installing Kubernetes with Azure Container Service

- Microsoft Azure offers a hosted Kubernetes-as-a-Service as part of the Azure Container Service
- The easiest way to get started with Azure Container Service is to use the built-in Azure Cloud Shell in the Azure portal
- You can activate the shell by clicking the shell icon
- Once you have the shell up and working, you can run:
`az group create --name=kuar --location=westus`
- Once the resource group is created, you can create a cluster using:
`az acs create --orchestrator-type=kubernetes --resource-group=kuar --name=kuar-cluster`
- Once the cluster is created, you can get credentials for the cluster with:
`az acs kubernetes get-credentials --resource-group=kuar --name=kuar-cluster`
- If you don't already have the kubectl tool installed, you can install it using:
`az acs kubernetes install-cli`

Installing Kubernetes on Amazon AWS

- Amazon Web Services (AWS) recently introduced a managed Kubernetes service called EKS
- It's still under preview mode
- At the moment Kubernetes can be installed on AWS as explained in the Kubernetes documentation either using conjure-up, Kubernetes Operations (kops), CoreOS Tectonic or kube-aws

Installing Kubernetes Locally Using minikube

- If you need a local development experience, or you don't want to pay for cloud resources, you can install a simple single-node cluster using minikube
- While minikube is a good simulation of a Kubernetes cluster, it is really intended for local development, learning, and experimentation
- Because it only runs in a VM on a single node, it doesn't provide the reliability of a distributed Kubernetes cluster
- You can find the minikube tool on GitHub
- There are binaries for Linux, macOS, and Windows that you can download
- Once you have the minikube tool installed you can create a local cluster using:
`minikube start`
- This will create a local VM, provision Kubernetes, and create a local kubectl configuration that points to that cluster
- When you are done with your cluster, you can stop the VM with:
`minikube stop`
- If you want to remove the cluster, you can run:
`minikube delete`

Installing Kubernetes on Bare Metal

- We will use Amazon AWS cloud for native installation and setup of Kubernetes
- Exercise 1: Install Kubernetes on Bare Metal

Kubernetes - Kubectl

- Kubectl is the command line utility to interact with Kubernetes API
- **Kubectl** controls the Kubernetes Cluster
- It is one of the key components of Kubernetes which runs on the workstation on any machine when the setup is done
- It has the capability to manage the nodes in the cluster.
- **Kubectl** commands are used to interact and manage Kubernetes objects and the cluster

Kubernetes Images

- Kubernetes (Docker) images are the key building blocks of Containerized Infrastructure
- Each container in a pod has its Docker image running inside it
- When we are configuring a pod, the image property in the configuration file has the same syntax as the Docker command does
- The configuration file has a field to define the image name, which we are planning to pull from the registry

Kubernetes Images

- In order to pull the image and create a container, we will run the following command

```
kubectl create -f Testing_for_Image_pull
```
- Once we fetch the log, we will get the output as successful

```
kubectl log Testing_for_Image_pull
```
- The above command will produce an output of success or we will get an output as failure

Kubernetes Jobs

- The main function of a job is to create one or more pod and tracks about the success of pods
- Jobs ensure that the specified number of pods are completed successfully
- When a specified number of successful run of pods is completed, then the job is considered complete

Creating a Job

- We will create the job using the following command with yaml which is saved with the name **py.yaml**

```
kubectl create -f py.yaml
```

- The above command will create a job
- If you want to check the status of a job, use the following command

```
kubectl describe jobs/py
```

Scheduled Jobs

- Scheduled job in Kubernetes uses **Cronetes**, which takes Kubernetes job and launches them in Kubernetes cluster
- Scheduling a job will run a pod at a specified point of time
- A parodic job is created for it which invokes itself automatically
- We will use the same yaml which we used to create the job and make it a scheduled job
- This scheduled job concept is useful when we are trying to build and run a set of tasks at a specified point of time and then complete the process

Kubernetes – Labels & Selectors

- **Labels**

- Labels are key-value pairs which are attached to pods, replication controller and services
- They are used as identifying attributes for objects such as pods and replication controller.
- They can be added to an object at creation time and can be added or modified at the run time

- **Selectors**

- Labels do not provide uniqueness
- In general, we can say many objects can carry the same labels
- Labels selector are core grouping primitive in Kubernetes
- They are used by the users to select a set of objects.
- Kubernetes API currently supports two type of selectors –
 - Equality-based selectors
 - Set-based selectors

Namespace

- Namespace provides an additional qualification to a resource name
- This is helpful when multiple teams are using the same cluster and there is a potential of name collision
- It can be as a virtual wall between multiple clusters
- Following are some of the important functionalities of a Namespace in Kubernetes –
 - Namespaces help pod-to-pod communication using the same namespace
 - Namespaces are virtual clusters that can sit on top of the same physical cluster
 - They provide logical separation between the teams and their environments

Kubernetes - Node

- A node is a working machine in Kubernetes cluster which is also known as a minion
- They are working units which can be physical, VM, or a cloud instance
- Each node has all the required configuration required to run a pod on it such as the proxy service and kubelet service along with the Docker, which is used to run the Docker containers on the pod created on the node
- They are not created by Kubernetes but they are created externally either by the cloud service provider or the Kubernetes cluster manager on physical or VM machines
- The key component of Kubernetes to handle multiple nodes is the controller manager, which runs multiple kind of controllers to manage nodes
- To manage nodes, Kubernetes creates an object of kind node which will validate that the object which is created is a valid node

Node Controller

- They are the collection of services which run in the Kubernetes master and continuously monitor the node in the cluster on the basis of metadata.name
- If all the required services are running, then the node is validated and a newly created pod will be assigned to that node by the controller
- If it is not valid, then the master will not assign any pod to it and will wait until it becomes valid.
- Kubernetes master registers the node automatically, if `--register-node` flag is true
 - `--register-node = true`
- If the cluster administrator wants to manage it manually then it could be done by turning the flag off –
 - `--register-node = false`

Kubernetes - Service

- A service can be defined as a logical set of pods
- It can be defined as an abstraction on the top of the pod which provides a single IP address and DNS name by which pods can be accessed
- With Service, it is very easy to manage load balancing configuration
- It helps pods to scale very easily.
- A service is a REST object in Kubernetes whose definition can be posted to Kubernetes apiServer on the Kubernetes master to create a new instance

Kubernetes – Replication Controller

- Replication Controller is one of the key features of Kubernetes, which is responsible for managing the pod lifecycle
- It is responsible for making sure that the specified number of pod replicas are running at any point of time
- It is used in time when one wants to make sure that the specified number of pod or at least one pod is running
- It has the capability to bring up or down the specified no of pod
- It is a best practice to use the replication controller to manage the pod life cycle rather than creating a pod again and again

Kubernetes – Replica Sets

- Replica Set ensures how many replica of pod should be running
- It can be considered as a replacement of replication controller
- The key difference between the replica set and the replication controller is, the replication controller only supports equality-based selector whereas the replica set supports set-based selector

Kubernetes – Deployments

- Deployments are upgraded and higher version of replication controller
- They manage the deployment of replica sets which is also an upgraded version of the replication controller.
- They have the capability to update the replica set and are also capable of rolling back to the previous version
- They provide many updated features of **matchLabels** and **selectors**
- We have got a new controller in the Kubernetes master called the deployment controller which makes it happen
- It has the capability to change the deployment midway

Kubernetes – Autoscaling

- **Autoscaling** is one of the key features in Kubernetes cluster
- It is a feature in which the cluster is capable of increasing the number of nodes as the demand for service response increases and decrease the number of nodes as the requirement decreases
- This feature of auto scaling is currently supported in Google Cloud Engine (GCE) and Google Container Engine (GKE) and will start with AWS pretty soon
- In order to set up scalable infrastructure in GCE, we need to first have an active GCE project with features of Google cloud monitoring, google cloud logging, and stackdriver enabled



THANK YOU